

```
# =====
# ⚡ AIM: Build Image Classification Model (CIFAR-10)
# =====

# -----
# a. Loading & Preprocessing
# -----

import tensorflow as tf

from tensorflow.keras import datasets, layers, models

import matplotlib.pyplot as plt

import numpy as np

import seaborn as sns

# Load CIFAR-10 dataset

(X_train, y_train), (X_test, y_test) = datasets.cifar10.load_data()

y_train = y_train.flatten()

y_test = y_test.flatten()

# Normalize pixel values (0–1)

X_train, X_test = X_train / 255.0, X_test / 255.0

# Class names for CIFAR-10

class_names = ['Airplane','Automobile','Bird','Cat','Deer',
'Dog','Frog','Horse','Ship','Truck']

print("✅ Training samples:", X_train.shape)
```

```
print("✓ Test samples:", X_test.shape)

# Show sample images
plt.figure(figsize=(8,8))
for i in range(9):
    plt.subplot(3,3,i+1)
    plt.imshow(X_train[i])
    plt.title(class_names[y_train[i]])
    plt.axis('off')
plt.suptitle("Sample CIFAR-10 Images", fontsize=14)
plt.show()

# -----
# b. Define Model Architecture
# -----
model = models.Sequential([
    layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64, (3,3), activation='relu'),
    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64, (3,3), activation='relu'),
    layers.Flatten(),

    layers.Dense(64, activation='relu'),
```

```
    layers.Dense(10, activation='softmax') # 10 output classes  
])
```

```
model.summary()
```

```
# -----
```

```
# c. Train the Model
```

```
# -----
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
history = model.fit(X_train, y_train, epochs=2,  
                     validation_data=(X_test, y_test),  
                     batch_size=64, verbose=1)
```

```
# -----
```

```
# d. Evaluate & Analyze Performance
```

```
# -----
```

```
test_loss, test_acc = model.evaluate(X_test, y_test, verbose=2)  
print(f"\n ✅ Test Accuracy: {test_acc*100:.2f}%")
```

```
# Plot accuracy & loss curves
```

```
plt.figure(figsize=(12,5))
```

```
plt.subplot(1,2,1)
```

```
plt.plot(history.history['accuracy'], label='Train Accuracy')
```

```
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')

plt.legend(); plt.title("Model Accuracy"); plt.xlabel("Epochs"); plt.ylabel("Accuracy")

plt.subplot(1,2,2)

plt.plot(history.history['loss'], label='Train Loss')

plt.plot(history.history['val_loss'], label='Validation Loss')

plt.legend(); plt.title("Model Loss"); plt.xlabel("Epochs"); plt.ylabel("Loss")

plt.show()

# ----

# 📈 Histogram of Predictions

# ----

y_pred = np.argmax(model.predict(X_test), axis=1)

print("Predicted Class:", class_names[y_pred[0]])

print("True Class:", class_names[y_test[0]])


plt.figure(figsize=(8,5))

sns.histplot(y_pred, bins=10, kde=False)

plt.title("Histogram of Predicted Classes")

plt.xlabel("Predicted Class (0–9)")

plt.ylabel("Frequency")

plt.show()

# ----

# 💡 Anomaly Detection (Misclassified Images)
```

```
# -----
misclassified_idx = np.where(y_pred != y_test)[0]
print(f"⚠️ Misclassified samples: {len(misclassified_idx)}")

# Display 9 misclassified images
plt.figure(figsize=(9,9))

for i, idx in enumerate(misclassified_idx[:9]):
    plt.subplot(3,3,i+1)
    plt.imshow(X_test[idx])
    plt.title(f"True: {class_names[y_test[idx]]}\nPred: {class_names[y_pred[idx]]}")
    plt.axis('off')

plt.suptitle("Anomalous (Misclassified) Images", fontsize=14)
plt.show()
```